

ICP 7 REPORT

```
# Mount Google Drive
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
import numpy as np
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPooling2D
from tensorflow.keras.constraints import MaxNorm
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.utils import to_categorical

# fix random seed for reproducibility
seed = 7
np.random.seed(seed)

# load data
(X_train, y_train), (X_test, y_test) = cifar10.load_data()

# normalize inputs from 0-255 to 0.0-1.0
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
X_train /= 255.0
X_test /= 255.0

# one hot encode outputs
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
num_classes = y_test.shape[1]

# Create the modified model
model = Sequential()

# Convolutional input layer, 32 feature maps with a size of 3x3, and a rectifier activation function.
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
```

```
# Convolutional layer, 32 feature maps with a size of 3x3 and a rectifier activation function.
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))

# Convolutional layer, 64 feature maps with a size of 3x3 and a rectifier activation function.
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))

# Convolutional layer, 64 feature maps with a size of 3x3 and a rectifier activation function.
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))

# Convolutional layer, 128 feature maps with a size of 3x3 and a rectifier activation function.
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))

# Convolutional layer, 128 feature maps with a size of 3x3 and a rectifier activation function.
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dropout(0.2))

# Fully connected layer with 1024 units and a rectifier activation function.
model.add(Dense(1024, activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))

# Fully connected layer with 512 units and a rectifier activation function.
model.add(Dense(512, activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))

# Fully connected output layer with 10 units and a Softmax activation function.
model.add(Dense(num_classes, activation='softmax'))

# Compile model
epochs = 25
lr_rate = 0.01
decay = lr_rate/epochs
sgd = SGD(learning_rate=lr_rate, momentum=0.9, nesterov=False)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
```

```
print(model.summary())

# Train the model
# Uncomment the line below to train in Colab
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)

# Evaluate the model
# Uncomment the lines below to evaluate in Colab
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 [=====] - 83s 0us/step
Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
dropout (Dropout)	(None, 32, 32, 32)	0
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
dropout_1 (Dropout)	(None, 16, 16, 64)	0
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73856
dropout_2 (Dropout)	(None, 8, 8, 128)	0
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0

```
flatten (Flatten) (None, 2048) 0
dropout_3 (Dropout) (None, 2048) 0
dense (Dense) (None, 1024) 2098176
dropout_4 (Dropout) (None, 1024) 0
dense_1 (Dense) (None, 512) 524800
dropout_5 (Dropout) (None, 512) 0
dense_2 (Dense) (None, 10) 5130

Total params: 2915114 (11.12 MB)
Trainable params: 2915114 (11.12 MB)
Non-trainable params: 0 (0.00 Byte)
```

Epoch	1/25	2/25	3/25	4/25	5/25	6/25	7/25	8/25	9/25	10/25	11/25	12/25	13/25
1563/1563	[=====] - 27s 10ms/step - loss: 1.9054 - accuracy: 0.2915 - val_loss: 1.6445 - val_accuracy: 0.3984	[=====] - 14s 9ms/step - loss: 1.4360 - accuracy: 0.4771 - val_loss: 1.2446 - val_accuracy: 0.5511	[=====] - 14s 9ms/step - loss: 1.2246 - accuracy: 0.5606 - val_loss: 1.1055 - val_accuracy: 0.6133	[=====] - 14s 9ms/step - loss: 1.0571 - accuracy: 0.6244 - val_loss: 0.9660 - val_accuracy: 0.6651	[=====] - 14s 9ms/step - loss: 0.9412 - accuracy: 0.6679 - val_loss: 0.8839 - val_accuracy: 0.6935	[=====] - 14s 9ms/step - loss: 0.8579 - accuracy: 0.6986 - val_loss: 0.8788 - val_accuracy: 0.6966	[=====] - 14s 9ms/step - loss: 0.7914 - accuracy: 0.7232 - val_loss: 0.7880 - val_accuracy: 0.7240	[=====] - 14s 9ms/step - loss: 0.7383 - accuracy: 0.7408 - val_loss: 0.7482 - val_accuracy: 0.7427	[=====] - 13s 8ms/step - loss: 0.6952 - accuracy: 0.7578 - val_loss: 0.7716 - val_accuracy: 0.7352	[=====] - 14s 9ms/step - loss: 0.6623 - accuracy: 0.7661 - val_loss: 0.7041 - val_accuracy: 0.7580	[=====] - 13s 8ms/step - loss: 0.6287 - accuracy: 0.7793 - val_loss: 0.7048 - val_accuracy: 0.7519	[=====] - 13s 8ms/step - loss: 0.6101 - accuracy: 0.7879 - val_loss: 0.7356 - val_accuracy: 0.7466	[=====] - 13s 8ms/step - loss: 0.6101 - accuracy: 0.7879 - val_loss: 0.7356 - val_accuracy: 0.7466

```

1563/1563 [=====] - 13s 8ms/step - loss: 0.6101 - accuracy: 0.7879 - val_loss: 0.7356 - val_accuracy: 0.7466
Epoch 13/25
1563/1563 [=====] - 14s 9ms/step - loss: 0.5920 - accuracy: 0.7928 - val_loss: 0.6814 - val_accuracy: 0.7692
Epoch 14/25
1563/1563 [=====] - 14s 9ms/step - loss: 0.5724 - accuracy: 0.8011 - val_loss: 0.6911 - val_accuracy: 0.7657
Epoch 15/25
1563/1563 [=====] - 14s 9ms/step - loss: 0.5594 - accuracy: 0.8031 - val_loss: 0.6881 - val_accuracy: 0.7692
Epoch 16/25
1563/1563 [=====] - 13s 9ms/step - loss: 0.5547 - accuracy: 0.8042 - val_loss: 0.6823 - val_accuracy: 0.7722
Epoch 17/25
1563/1563 [=====] - 14s 9ms/step - loss: 0.5400 - accuracy: 0.8124 - val_loss: 0.7263 - val_accuracy: 0.7579
Epoch 18/25
1563/1563 [=====] - 14s 9ms/step - loss: 0.5363 - accuracy: 0.8143 - val_loss: 0.7219 - val_accuracy: 0.7610
Epoch 19/25
1563/1563 [=====] - 14s 9ms/step - loss: 0.5181 - accuracy: 0.8221 - val_loss: 0.7157 - val_accuracy: 0.7603
Epoch 20/25
1563/1563 [=====] - 13s 8ms/step - loss: 0.5209 - accuracy: 0.8174 - val_loss: 0.7396 - val_accuracy: 0.7532
Epoch 21/25
1563/1563 [=====] - 13s 8ms/step - loss: 0.5135 - accuracy: 0.8244 - val_loss: 0.6939 - val_accuracy: 0.7684
Epoch 22/25
1563/1563 [=====] - 14s 9ms/step - loss: 0.5226 - accuracy: 0.8195 - val_loss: 0.7000 - val_accuracy: 0.7685
Epoch 23/25
1563/1563 [=====] - 14s 9ms/step - loss: 0.5253 - accuracy: 0.8195 - val_loss: 0.7049 - val_accuracy: 0.7665
Epoch 24/25
1563/1563 [=====] - 13s 8ms/step - loss: 0.5173 - accuracy: 0.8235 - val_loss: 0.8543 - val_accuracy: 0.7179
Epoch 25/25
1563/1563 [=====] - 14s 9ms/step - loss: 0.5404 - accuracy: 0.8175 - val_loss: 0.7220 - val_accuracy: 0.7672
Accuracy: 76.72%

```

```

import numpy as np

# Predict the first 4 images
predictions = model.predict(X_test[:4])

# Convert predictions from one-hot encoded to label indices
predicted_classes = np.argmax(predictions, axis=1)

# Convert actual labels from one-hot encoded to label indices
actual_classes = np.argmax(y_test[:4], axis=1)

# Print the results
for i in range(4):
    print(f"IMAGE {i+1}:")
    print(f"PREDICTED CLASS: {predicted_classes[i]}, ACTUAL CLASS: {actual_classes[i]}")
    if predicted_classes[i] == actual_classes[i]:
        print("PREDICTION IS CORRECT!")
    else:
        print("PREDICTION IS INCORRECT!")
    print("-----")

```

```

1/1 [=====] - 0s 301ms/step
IMAGE 1:
PREDICTED CLASS: 3, ACTUAL CLASS: 3
PREDICTION IS CORRECT!
-----
IMAGE 2:
PREDICTED CLASS: 8, ACTUAL CLASS: 8
PREDICTION IS CORRECT!
-----
IMAGE 3:
PREDICTED CLASS: 8, ACTUAL CLASS: 8
PREDICTION IS CORRECT!
-----
IMAGE 4:
PREDICTED CLASS: 0, ACTUAL CLASS: 0
PREDICTION IS CORRECT!
-----

```

```

import matplotlib.pyplot as plt


# Plotting the Loss
plt.figure(figsize=(12, 4))

plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Loss vs. Epochs')
plt.xlabel('EPOCHS')
plt.ylabel('LOSS')
plt.legend()

# Plotting the Accuracy
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy vs. Epochs')
plt.xlabel('EPOCHS')
plt.ylabel('ACCURACY')
plt.legend()

plt.tight_layout()
plt.show()

```



GITHUB Repo: <https://github.com/sxk7912/Bigdata>